

Analysis of Time Spent in Shelter

Group 20

```
library(MASS)
library(stats)
library(ggplot2)

# Read the CSV file into a dataframe
df <- read.csv("dataset20.csv")

# Ensure the 'month' column is numeric
df$month <- as.numeric(df$month)

# Categorize months into seasons
df$season <- cut(df$month, breaks = c(2, 5, 8, 11, 12),
                labels = c('Spring', 'Summer', 'Autumn', 'Winter'),
                include.lowest = TRUE)

# Correct the classification of December, January, and February as Winter
df$season[df$month %in% c(12, 1, 2)] <- 'Winter'

# Convert 'season' to a factor and set the order
df$season <- factor(df$season, levels = c("Spring", "Summer", "Autumn", "Winter"))
```

Spring is the breeding season for many animals, which may lead to a surge in the number of stray animals, putting shelters under immense pressure. As resources become strained and workloads increase, staff efficiency may decline, resulting in longer stays for animals in shelters.

Summer is a peak travel season, and as people are away from home, the demand for pet adoption decreases. Additionally, the hot and humid environment increases the risk of diseases, further prolonging the stay of animals in shelters.

In winter, the holiday season, including Christmas and New Year, may lead to an adoption surge, reducing the time animals spend in shelters.

Given the impact of these seasonal factors on animal sheltering and adoption trends, we have decided to incorporate seasonal effects as an explanatory variable

```
# Fit a negative binomial regression model
full_model <- glm.nb(time_at_shelter ~ animal_type + intake_type + chip_status + season + year)
```

Outcome_type is an outcome variable rather than a factor influencing the length of stay in the shelter, and therefore should not be used as an explanatory variable. In this study, we use animal_type, chip_status, intake_type, season, and year as explanatory variables, while time_at_shelter serves as the outcome variable to construct a negative binomial regression model.

```
# Perform backward stepwise selection to simplify the model
selected_model1 <- step(full_model, direction = "backward", trace = TRUE)
```

Start: AIC=7640.67

```
time_at_shelter ~ animal_type + intake_type + chip_status + season +
  year
```

	Df	Deviance	AIC
- year	1	1638.2	7639.1
<none>		1637.7	7640.7
- season	3	1644.5	7641.5
- animal_type	2	1644.2	7643.2
- chip_status	2	1661.0	7660.0
- intake_type	2	1697.3	7696.2

Step: AIC=7639.12

```
time_at_shelter ~ animal_type + intake_type + chip_status + season
```

	Df	Deviance	AIC
<none>		1637.6	7639.1
- season	3	1645.0	7640.4
- animal_type	2	1644.2	7641.7
- chip_status	2	1660.9	7658.3
- intake_type	2	1697.3	7694.8

To select the most appropriate explanatory variables, we employ a backward stepwise regression approach using the Akaike Information Criterion (AIC) as the evaluation standard. Specifically, we start with a full model that includes all candidate explanatory variables and iteratively remove variables with lower contributions to the model until we identify the optimal

model with the lowest AIC. The results indicate that removing the season variable leads to a lower AIC value for the model. Therefore, we exclude the season variable to improve model fit.

```
# Display the summary of the selected model
summary(selected_model1)
```

Call:

```
glm.nb(formula = time_at_shelter ~ animal_type + intake_type +
        chip_status + season, data = df, init.theta = 0.9091538376,
        link = log)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	1.88927	0.16682	11.325	< 2e-16 ***
animal_typeDOG	0.20526	0.07838	2.619	0.008826 **
animal_typeWILDLIFE	0.36788	0.57199	0.643	0.520125
intake_typeOWNER SURRENDER	-0.94831	0.13746	-6.899	5.24e-12 ***
intake_typeSTRAY	-0.57829	0.13094	-4.416	1.00e-05 ***
chip_statusSCAN NO CHIP	0.17055	0.08119	2.101	0.035662 *
chip_statusUNABLE TO SCAN	-0.66931	0.18734	-3.573	0.000353 ***
seasonSummer	0.14008	0.08112	1.727	0.084210 .
seasonAutumn	0.10065	0.08712	1.155	0.247977
seasonWinter	0.24006	0.09077	2.645	0.008174 **

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(0.9092) family taken to be 1)

Null deviance: 1737.3 on 1404 degrees of freedom
Residual deviance: 1637.6 on 1395 degrees of freedom
AIC: 7641.1

Number of Fisher Scoring iterations: 1

Theta: 0.9092
Std. Err.: 0.0459

2 x log-likelihood: -7619.1160

After removing the season variable, we reconstructed the negative binomial regression model. The results indicate that at a 95% confidence level, the p-value of the year variable is greater

than 0.05, suggesting that it is not statistically significant. Additionally, based on the AIC evaluation, the model with the year variable has an AIC of 8353.1, while the model without it has an AIC of 8353.6, indicating that removing the year variable has a minimal impact on the model. Therefore, we decide to exclude the year variable from the explanatory variables.

```
# Fit a negative binomial regression model with a reduced set of predictors
selected_model2 <- glm.nb(time_at_shelter ~ animal_type + intake_type + chip_status, data = df)

# Display the summary of the model
summary(selected_model2)
```

Call:

```
glm.nb(formula = time_at_shelter ~ animal_type + intake_type +
        chip_status, data = df, init.theta = 0.9022135578, link = log)
```

Coefficients:

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	2.01893	0.15799	12.778	< 2e-16 ***
animal_typeDOG	0.20473	0.07799	2.625	0.008659 **
animal_typeWILDLIFE	0.37388	0.57373	0.652	0.514620
intake_typeOWNER SURRENDER	-0.93785	0.13783	-6.804	1.01e-11 ***
intake_typeSTRAY	-0.57653	0.13135	-4.389	1.14e-05 ***
chip_statusSCAN NO CHIP	0.14810	0.08121	1.824	0.068190 .
chip_statusUNABLE TO SCAN	-0.66838	0.18756	-3.564	0.000366 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(0.9022) family taken to be 1)

Null deviance: 1728.7 on 1404 degrees of freedom
 Residual deviance: 1636.9 on 1398 degrees of freedom
 AIC: 7642.4

Number of Fisher Scoring iterations: 1

Theta: 0.9022
 Std. Err.: 0.0454


2 x log-likelihood: -7626.4060

Finally, we select animal_type, chip_status, and intake_type as explanatory variables, with time_at_shelter as the outcome variable to construct a negative binomial regression model.

```
# Compute Pearson residuals from the model
df$residuals <- residuals(selected_model2, type = "pearson")

# Extract fitted values from the model
df$fitted_values <- fitted(selected_model2)


# Create scatter plot of fitted values vs. residuals
ggplot(df, aes(x = fitted_values, y = residuals)) +
  geom_point(alpha = 0.6) +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  theme_minimal() +
  labs(title = "Fitted Values vs. Pearson Residuals",
       x = "Fitted Values",
       y = "Pearson Residuals")
```



The plot shows some points with Pearson Residuals > 6 , indicating the possible presence of outliers in the data. These outliers may have a significant impact on the model, requiring further investigation to determine whether any adjustments or modifications to the model are necessary.

```
# Create a stem plot of Cook's Distance values
plot(cooks.distance(selected_model2), type="h",
     main="Cook's Distance", ylab="Cook's Distance")
```

```
# Add a red dashed horizontal line at Cook's Distance = 1
abline(h = 1, col = "red", lty = 2)
```



The highest Cook's Distance in the plot appears to be less than 0.1, which is far below 1, indicating that there are no particularly severe high-influence points. However, some points still exhibit a relatively large influence. Therefore, in the subsequent steps, we will identify these high-influence points and attempt to remove the outliers before refitting the model to assess their impact.

```
# Remove highly influential observations based on Cook's Distance
df_cleaned <- df[-which(cooks.distance(selected_model2) > 4 / nrow(df)), ]

# Refit the model using cleaned data
final_model <- glm.nb(time_at_shelter ~ animal_type + intake_type + chip_status, data = df_cleaned)

# Display summary of the final refined model
summary(final_model)
```

Call:

```
glm.nb(formula = time_at_shelter ~ animal_type + intake_type +
      chip_status, data = df_cleaned, init.theta = 0.9543940796,
```

```

link = log)

Coefficients:
                Estimate Std. Error z value Pr(>|z|)
(Intercept)      1.8929    0.1570  12.055  < 2e-16 ***
animal_typeDOG    0.2884    0.0789   3.656 0.000257 ***
animal_typeWILDLIFE 0.5395    0.5608   0.962 0.336048
intake_typeOWNER SURRENDER -1.0219    0.1373  -7.441 9.97e-14 ***
intake_typeSTRAY  -0.5597    0.1306  -4.284 1.84e-05 ***
chip_statusSCAN NO CHIP  0.1675    0.0807   2.075 0.037948 *
chip_statusUNABLE TO SCAN -1.3559    0.2380  -5.696 1.23e-08 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

(Dispersion parameter for Negative Binomial(0.9544) family taken to be 1)

Null deviance: 1740.2  on 1370  degrees of freedom
Residual deviance: 1599.6  on 1364  degrees of freedom
AIC: 7306.1

Number of Fisher Scoring iterations: 1

        Theta: 0.9544
      Std. Err.: 0.0501

2 x log-likelihood: -7290.0980

```

The results show that after removing the outliers, the refitted model has an AIC of 7642.4, which is 695.199 lower than the previous model's AIC, indicating an improvement in model fit.

When all categorical variables are set to their baseline categories (animal_type = cat, intake_type = confiscated, chip_status = scan chip), the model's log-predicted value is 2.01893. Thus, the estimated shelter stay for the baseline group (cats) is approximately 7.52 days.

The coefficient for animal_type = DOG is 0.20473, indicating that with other variables held constant, dogs stay longer in the shelter compared to the baseline category (cats), with an estimated increase of approximately $\exp(0.20473)$ 1.23 times. Additionally, animal_type = WILDLIFE is not statistically significant at the 95% confidence level, suggesting that the shelter stay duration for wildlife does not significantly differ from that of the baseline category (cats)

intake_type = OWNER SURRENDER has a significant impact on shelter stay duration ($p < 0.001$). With other variables held constant, animals surrendered by their owners stay in

the shelter for a shorter duration compared to the baseline category (confiscated), with a stay duration of $\exp(-0.93785) = 39\%$ of the baseline category. `intake_type = STRAY` also has a significant effect on shelter stay duration ($p < 0.001$). With other variables held constant, stray animals stay in the shelter for $\exp(-0.57653) = 56\%$ of the baseline category (confiscated).


`chip_status = SCAN NO CHIP` is not statistically significant at the 95% confidence level ($p = 0.068$), indicating that the shelter stay duration of animals without a chip does not significantly differ from that of the baseline category (SCAN CHIP). `chip_status = UNABLE TO SCAN` has a significant impact on shelter stay duration ($p = 0.0004$). With other variables held constant, animals whose chips cannot be scanned have a stay duration of only $\exp(-0.66838) = 0.51$ of the baseline category (SCAN CHIP), meaning their stay is approximately 51% of the baseline category.

1 Assessing model fit

```
# Compute Pearson residuals from the final model
df_cleaned$residuals <- residuals(final_model, type = "pearson")

# Extract fitted values from the final model
df_cleaned$fitted_values <- fitted(final_model)


# Create scatter plot of fitted values vs. residuals
ggplot(df_cleaned, aes(x = fitted_values, y = residuals)) +
  geom_point(alpha = 0.6) +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  theme_minimal() +
  labs(title = "Fitted Values vs. Pearson Residuals",
       x = "Fitted Values",
       y = "Pearson Residuals")
```

The Pearson residuals are distributed around 0 without any distinct U-shaped or V-shaped patterns, indicating that the model does not suffer from severe systematic bias and has a good overall fit.

```
# Create a stem plot of Cook's Distance values
plot(cooks.distance(final_model), type="h",
     main="Cook's Distance", ylab="Cook's Distance")

# Add a red dashed horizontal line at Cook's Distance = 1
abline(h = 1, col = "red", lty = 2)
```



Cook's Distance values are generally low, indicating that no single data point has an excessively large influence on the model, ensuring stable model fitting.

```
# Compute dispersion parameter  
deviance(final_model) / df.residual(final_model)
```

```
[1] 1.172697
```

The calculated value of 1.17089 is slightly greater than 1, indicating a mild degree of overdispersion in the data. However, overall, the model fits well and can still accurately describe the data distribution.

2 Conclusion

This study analyzed factors affecting the shelter stay duration of animals using data from a Dallas animal shelter. Through exploratory data analysis and statistical modeling, we found that intake type and animal type significantly influence shelter stay duration, while chip status has a limited impact.

The final negative binomial regression model included animal type, intake type, and chip status as explanatory variables. After removing outliers, the model's AIC decreased by 695.199, indicating an improved fit.

Key findings:

- Dogs stay **1.23 times longer** in the shelter than cats.
- Stray animals and owner-surrendered animals stay shorter than confiscated animals (56% and 39% of the baseline category, respectively).
- Chip status generally does not significantly impact stay duration, except for “unable to scan” cases, where animals had 51% of the baseline stay duration.

The Pearson residuals analysis showed no severe systematic bias, and Cook’s Distance analysis confirmed model stability. The slightly over dispersed data ($1.17089 > 1$) suggests a good model fit.